Go to next chapter >

2 Managing threads

This chapter covers

- Starting threads, and various ways of specifying code to run on a new thread
- Waiting for a thread to finish versus leaving it to run
- Uniquely identifying threads

OK, so you've decided to use concurrency for your application. In particular, you've decided to use multiple threads. What now? How do you launch these threads, how do you check that they've finished, and how do you keep tabs on them? The C++ Standard Library makes most thread-management tasks relatively easy, with just about everything managed through the `std::thread` object associated with a given thread, as you'll see. For those tasks that aren't so straightforward, the library provides the flexibility to build what you need from the basic building blocks.

In this chapter, I'll start by covering the basics: launching a thread, waiting for it to finish, or running it in the background. We'll then proceed to look at passing additional parameters to the thread function when it's launched and how to transfer ownership of a thread from one `std::thread` object to another. Finally, we'll look at choosing the number of threads to use and identifying particular threads.

2.1   Basic thread management

Every C++ program has at least one thread, which is started by the C++ runtime: the thread running `main()`. Your program can then launch additional threads that have another function as the entry point. These threads then run concurrently with each other and with the initial thread. Just as the program exits when the program returns from `main()`, when the specified entry point function returns, the thread exits. As you'll see, if you have a `std::thread` object for a thread, you can wait for it to finish; but first you have to start it, so let's look at launching threads.

2.1.1 Launching a thread

As you saw in chapter 1, threads are started by constructing a `std::thread` object that specifies the task to run on that thread. In the simplest case, that task is just a plain, ordinary `void` -returning function that takes no parameters. This function runs on its own thread until it returns, and then the thread stops. At the other extreme, the task could be a function object that takes additional parameters and performs a series of independent operations that are specified through some kind of messaging system while it's running, and the thread stops only when it's signaled to do so, again via some kind of messaging system. It doesn't matter what the thread is going to do or where it's launched from, but starting a thread using the C++ Thread Library always boils down to constructing a `std::thread` object:

```
1  void do_some_work();
2  std::thread my_thread(do_some_work);
```

This is just about as simple as it gets. Of course, you have to make sure that the `<thread>` header is included so the compiler can see the definition of the `std::thread` class. As with much of the C++ Standard Library, `std::thread` works with any callable type, so you can pass an instance of a class with a function call operator to the `std::thread` constructor instead:

```
1  class background_task
2  {
3  public:
4      void operator()() const
5      {
6          do_something();
7          do_something_else();
8      }
9  };
10  background_task f;
11  std::thread my_thread(f);
```

In this case, the supplied function object is copied into the storage belonging to the newly created thread of execution and invoked from there. It's therefore essential that the copy behave equivalently to the original, or the result may not be what's expected.

One thing to consider when passing a function object to the thread constructor is to avoid what is dubbed "C++'s most vexing parse." If you pass a temporary rather than a named variable, then the syntax can be the same as that of a function declaration, in which case the compiler interprets it as such, rather than an object definition. For example,

```
1  std::thread my_thread(background_task());
```

declares a function `my_thread` that takes a single parameter (of type pointer to a function taking no parameters and returning a `background_task` object) and returns a `std::thread` object, rather than launching a new thread. You can avoid this by naming your function object as shown previously, by using an extra set of parentheses, or by using the new uniform initialization syntax, for example:

```
1  std::thread my_thread((background_task()));       1
2  std::thread my_thread{background_task()};         2
```

In the first example #1, the extra parentheses prevent interpretation as a function declaration, thus allowing `my_thread` to be declared as a variable of type `std::thread`. The second example #2 uses the new uniform initialization syntax with braces rather than parentheses, and thus would also declare a variable.

One type of callable object that avoids this problem is a lambda expression. This is a new feature from C++11 which essentially allows you to write a local function, possibly capturing some local variables and avoiding the need of passing additional arguments (see section 2.2). For full details on lambda expressions, see appendix A, section A.5. The previous example can be written using a lambda expression as follows:

```
1  std::thread my_thread([]{
2      do_something();
3      do_something_else();
4  });
```

Once you've started your thread, you need to explicitly decide whether to wait for it to finish (by joining with it—see section 2.1.2) or leave it to run on its own (by detaching it—see section 2.1.3). If you don't decide before the `std::thread` object is destroyed, then your program is terminated (the `std::thread` destructor calls `std::terminate()`). It's therefore imperative that you ensure that the thread is correctly joined or detached, even in the presence of exceptions. See section 2.1.3 for a technique to handle this scenario. Note that you only have to make this decision before the `std::thread` object is destroyed— the thread itself may well have finished long before you join with it or detach it, and if you detach it, then if the thread is still running, it will continue to do so, and may continue running long after the `std::thread` object is destroyed; it will only stop running when it finally returns from the thread function.

If you don't wait for your thread to finish, then you need to ensure that the data accessed by the thread is valid until the thread has finished with it. This isn't a new problem—even in single-threaded code it is undefined behavior to access an object after it's been destroyed— but the use of threads provides an additional opportunity to encounter such lifetime issues.

One situation in which you can encounter such problems is when the thread function holds pointers or references to local variables and the thread hasn't finished when the function exits. The following listing shows an example of just such a scenario.

Listing 2.1 A function that returns while a thread still has access to local variables

1

```
 1  struct func
 2  {
 3      int& i;
 4      func(int& i_):i(i_){}
 5      void operator()()
 6      {
 7          for(unsigned j=0;j<1000000;++j)
 8          {
 9              do_something(i);
10          }
11      }
12  };
13  void oops()
14  {
15      int some_local_state=0;
16      func my_func(some_local_state);
17      std::thread my_thread(my_func);
18      my_thread.detach();
19  }
```

In this case, the new thread associated with `my_thread` will probably still be running when `oops` exits #3, because you've explicitly decided not to wait for it by calling `detach() #2`. If the thread is still running, then we have the scenario shown in figure 2.1: the next call to `do_something(i) #1` will access an already destroyed variable. This is just like normal single-threaded code—allowing a pointer or reference to a local variable to persist beyond the function exit is never a good idea—but it's easier to make the mistake with multithreaded code, because it isn't necessarily immediately apparent that this has happened.

Main thread	New thread
Construct `my_func` with reference to `some_local_state`	
Start new thread `my_thread`	
	Started
	Calls `func::operator()`

Detaches `my_thread`	Running `func::operator()`; may call `do_something` with reference to `some_local_state`
Destroys `some_local_state`	Still running
Exits `oops`	Still running `func::operator()`; may call `do_something` with reference to `some_local_state` => undefined behavior

Figure 2.1: Accessing a local variable with a detached thread after it has been destroyed

One common way to handle this scenario is to make the thread function self-contained and copy the data into the thread rather than sharing the data. If you use a callable object for your thread function, that object is itself copied into the thread, so the original object can be destroyed immediately. But you still need to be wary of objects containing pointers or references, such as that from listing 2.1. In particular, it's a bad idea to create a thread within a function that has access to the local variables in that function, unless the thread is guaranteed to finish before the function exits.

Alternatively, you can ensure that the thread has completed execution before the function exits by joining with the thread.

2.1.2 Waiting for a thread to complete

If you need to wait for a thread to complete, you can do this by calling `join()` on the associated `std::thread` instance. In the case of listing 2.1, replacing the call to `my_thread.detach()` before the closing brace of the function body with a call to `my_thread.join()` would therefore be sufficient to ensure that the thread was finished before the function was exited and thus before the local variables were destroyed. In this case, it would mean there was little point running the function on a separate thread, because the first thread wouldn't be doing anything useful in the meantime, but in real code the original thread would either have work to do itself or it would have launched several threads to do useful work before waiting for all of them to complete.

`join()` is simple and brute force—either you wait for a thread to finish or you don't. If you need more fine-grained control over waiting for a thread, such as to check whether a

thread is finished, or to wait only a certain period of time, then you have to use alternative mechanisms such as condition variables and futures, which we'll look at in chapter 4. The act of calling `join()` also cleans up any storage associated with the thread, so the `std::thread` object is no longer associated with the now-finished thread; it isn't associated with any thread. This means that you can call `join()` only once for a given thread; once you've called `join()`, the `std::thread` object is no longer joinable, and `joinable()` will return `false`.

2.1.3 Waiting in exceptional circumstances

As mentioned earlier, you need to ensure that you've called either `join()` or `detach()` before a `std::thread` object is destroyed. If you're detaching a thread, you can usually call `detach()` immediately after the thread has been started, so this isn't a problem. But if you're intending to wait for the thread, you need to pick carefully the place in the code where you call `join()`. This means that the call to `join()` is liable to be skipped if an exception is thrown after the thread has been started but before the call to `join()`.

To avoid your application being terminated when an exception is thrown, you therefore need to make a decision on what to do in this case. In general, if you were intending to call `join()` in the non-exceptional case, you also need to call `join()` in the presence of an exception to avoid accidental lifetime problems. The next listing shows some simple code that does just that.

Listing 2.2 Waiting for a thread to finish

```
1   struct func;                                            A
2   void f()
3   {
4       int some_local_state=0;
5       func my_func(some_local_state);
6       std::thread t(my_func);
7       try
8       {
9           do_something_in_current_thread();
10      }
11      catch(...)
12      {
13          t.join();                                       1
14          throw;
15      }
16      t.join();                                           2
17  }
```

The code in listing 2.2 uses a $\boxed{\texttt{try}}$ / $\boxed{\texttt{catch}}$ block to ensure that a thread with access to local state is finished before the function exits, whether the function exits normally #2 or by an exception #1. The use of $\boxed{\texttt{try}}$ / $\boxed{\texttt{catch}}$ blocks is verbose, and it's easy to get the scope slightly wrong, so this isn't an ideal scenario. If it's important to ensure that the thread must complete before the function exits—whether because it has a reference to other local variables or for any other reason—then it's important to ensure this is the case for all possible exit paths, whether normal or exceptional, and it's desirable to provide a simple, concise mechanism for doing so.

One way of doing this is to use the standard Resource Acquisition Is Initialization (RAII) idiom and provide a class that does the $\boxed{\texttt{join()}}$ in its destructor, as in the following listing. See how it simplifies the function $\boxed{\texttt{f()}}$.

Listing 2.3 Using RAII to wait for a thread to complete

```
 1  class thread_guard
 2  {
 3      std::thread& t;
 4  public:
 5      explicit thread_guard(std::thread& t_):
 6          t(t_)
 7      {}
 8      ~thread_guard()
 9      {
10          if(t.joinable())                                        1
11          {
12              t.join();                                           2
13          }
14      }
15      thread_guard(thread_guard const&)=delete;                   3
16      thread_guard& operator=(thread_guard const&)=delete;
17  };
18  struct func;                                                    A
19  void f()
20  {
21      int some_local_state=0;
22      func my_func(some_local_state);
23      std::thread t(my_func);
24      thread_guard g(t);
25      do_something_in_current_thread();
26  }                                                               4
```

When the execution of the current thread reaches the end of `f #4`, the local objects are destroyed in reverse order of construction. Consequently, the `thread_guard` object `g` is destroyed first, and the thread is joined with in the destructor #2. This even happens if the function exits because `do_something_in_current_thread` throws an exception.

The destructor of `thread_guard` in listing 2.3 first tests to see if the `std::thread` object is `joinable() #1` before calling `join() #2`. This is important, because `join()` can be called only once for a given thread of execution, so it would therefore be a mistake to do so if the thread had already been joined.

The copy constructor and copy-assignment operator are marked `=delete #3` to ensure that they're not automatically provided by the compiler. Copying or assigning such an object would be dangerous, because it might then outlive the scope of the thread it was joining. By declaring them as deleted, any attempt to copy a `thread_ guard` object will generate a compilation error. See appendix A, section A.2, for more about deleted functions.

If you don't need to wait for a thread to finish, you can avoid this exception-safety issue by detaching it. This breaks the association of the thread with the `std::thread` object and ensures that `std::terminate()` won't be called when the `std::thread` object is destroyed, even though the thread is still running in the background.

2.1.4 Running threads in the background

Calling `detach()` on a `std::thread` object leaves the thread to run in the background, with no direct means of communicating with it. It's no longer possible to wait for that thread to complete; if a thread becomes detached, it isn't possible to obtain a `std::thread` object that references it, so it can no longer be joined. Detached threads truly run in the background; ownership and control are passed over to the C++ Runtime Library, which ensures that the resources associated with the thread are correctly reclaimed when the thread exits.

Detached threads are often called daemon threads after the UNIX concept of a daemon process that runs in the background without any explicit user interface. Such threads are typically long-running; they may well run for almost the entire lifetime of the application, performing a background task such as monitoring the filesystem, clearing unused entries out of object caches, or optimizing data structures. At the other extreme, it may make sense to use a detached thread where there's another mechanism for identifying when the thread has completed or where the thread is used for a "fire and forget" task.

As you've already seen in section 2.1.2, you detach a thread by calling the `detach()` member function of the `std::thread` object. After the call completes, the `std::thread` object is no longer associated with the actual thread of execution and is therefore no longer joinable:

```
1  std::thread t(do_background_work);
2  t.detach();
3  assert(!t.joinable());
```

In order to detach the thread from a `std::thread` object, there must be a thread to detach: you can't call `detach()` on a `std::thread` object with no associated thread of execution. This is exactly the same requirement as for `join()`, and you can check it in exactly the same way—you can only call `t.detach()` for a `std::thread` object `t` when `t.joinable()` returns `true`.

Consider an application such as a word processor that can edit multiple documents at once. There are many ways to handle this, both at the UI level and internally. One way that seems to be increasingly common at the moment is to have multiple independent top-level windows, one for each document being edited. Although these windows appear to be completely independent, each with its own menus and so forth, they're running within the same instance of the application. One way to handle this internally is to run each document-editing window in its own thread; each thread runs the same code but with different data relating to the document being edited and the corresponding window properties. Opening a new document therefore requires starting a new thread. The thread handling the request isn't going to care about waiting for that other thread to finish, because it's working on an unrelated document, so this makes it a prime candidate for running a detached thread.

The following listing shows a simple code outline for this approach.

MEAP

Listing 2.4 Detaching a thread to handle other documents

```
 1  void edit_document(std::string const& filename)
 2  {
 3      open_document_and_display_gui(filename);
 4      while(!done_editing())
 5      {
 6          user_command cmd=get_user_input();
 7          if(cmd.type==open_new_document)
 8          {
 9              std::string const new_name=get_filename_from_user();
10              std::thread t(edit_document,new_name);
11              t.detach();
12          }
13          else
14          {
15              process_user_input(cmd);
16          }
17      }
18  }
```

If the user chooses to open a new document, you prompt them for the document to open, start a new thread to open that document #1, and then detach it #2. Because the new thread is doing the same operation as the current thread but on a different file, you can reuse the same function (`edit_document`) with the newly chosen filename as the supplied argument.

This example also shows a case where it's helpful to pass arguments to the function used to start a thread: rather than just passing the name of the function to the `std::thread` constructor #1, you also pass in the filename parameter. Although other mechanisms could be used to do this, such as using a function object with member data instead of an ordinary function with parameters, the Thread Library provides you with an easy way of doing it.

2.2   Passing arguments to a thread function

As shown in listing 2.4, passing arguments to the callable object or function is fundamentally as simple as passing additional arguments to the `std::thread` constructor. But it's important to bear in mind that by default the arguments are copied into internal storage, where they can be accessed by the newly created thread of execution, and then passed to the callable object or function as rvalues as if they were temporaries. This is done even if the corresponding parameter in the function is expecting a reference. Here's a simple example:

```
1  void f(int i,std::string const& s);
2  std::thread t(f,3,"hello");
```

This creates a new thread of execution associated with `t`, which calls `f(3,"hello")`. Note that even though `f` takes a `std::string` as the second parameter, the string literal is passed as a `char` `const*` and converted to a `std::string` only in the context of the new thread. This is particularly important when the argument supplied is a pointer to an automatic variable, as follows:

```
1  void f(int i,std::string const& s);
2  void oops(int some_param)
3  {
4      char buffer[1024];                       1
5      sprintf(buffer, "%i",some_param);
6      std::thread t(f,3,buffer);               2
7      t.detach();
8  }
```

In this case, it's the pointer to the local variable `buffer #1` that's passed through to the new thread #2, and there's a significant chance that the function `oops` will exit before the buffer has been converted to a `std::string` on the new thread, thus leading to undefined behavior. The solution is to cast to `std::string` before passing the buffer to the `std::thread` constructor:

```
1  void f(int i,std::string const& s);
2  void not_oops(int some_param)
3  {
4      char buffer[1024];
5      sprintf(buffer,"%i",some_param);
6      std::thread t(f,3,std::string(buffer));      A
7      t.detach();
8  }
```

In this case, the problem is that you were relying on the implicit conversion of the pointer to the buffer into the `std::string` object expected as a function parameter, but this conversion happens too late because the `std::thread` constructor copies the supplied values as is, without converting to the expected argument type.

It's not possible to get the reverse scenario: the object is copied, and what you wanted was a non- `const` reference, since this won't compile. You might try and do this if the thread is updating a data structure that's passed in by reference, for example:

```
1  void update_data_for_widget(widget_id w,widget_data& data);         1
2  void oops_again(widget_id w)
3  {
4      widget_data data;
5      std::thread t(update_data_for_widget,w,data);                   2
6      display_status();
7      t.join();
8      process_widget_data(data);
9  }
```

Although `update_data_for_widget #1` expects the second parameter to be passed by reference, the `std::thread` constructor #2 doesn't know that; it's oblivious to the types of the arguments expected by the function and blindly copies the supplied values. However, the internal code passes copied arguments as rvalues in order to work with move-only types, and will thus try to call update_data_for_widget with an rvalue. This will therefore fail to compile as you can't pass an rvalue to a function expecting a non- `const` reference. For those of you familiar with `std::bind`, the solution will be readily apparent: you need to wrap the arguments that really need to be references in `std::ref`. In this case, if you change the thread invocation to

```
1  std::thread t(update_data_for_widget,w,std::ref(data));
```

and then `update_data_for_widget` will be correctly passed a reference to `data` rather than a temporary copy of `data`, and the code will now compile successfully.

If you're familiar with `std::bind`, the parameter-passing semantics will be unsurprising, because both the operation of the `std::thread` constructor and the operation of `std::bind` are defined in terms of the same mechanism. This means that, for example, you can pass a member function pointer as the function, provided you supply a suitable object pointer as the first argument:

```
1  class X
2  {
3  public:
4      void do_lengthy_work();
5  };
6  X my_x;
7  std::thread t(&X::do_lengthy_work,&my_x);                              1
```

This code will invoke `my_x.do_lengthy_work()` on the new thread, because the address of `my_x` is supplied as the object pointer #1. You can also supply arguments to such a member function call: the third argument to the `std::thread` constructor will be the first argument to the member function and so forth.

Another interesting scenario for supplying arguments is where the arguments can't be copied but can only be moved: the data held within one object is transferred over to another, leaving the original object "empty." An example of such a type is `std::unique_ptr`, which provides automatic memory management for dynamically allocated objects. Only one `std::unique_ptr` instance can point to a given object at a time, and when that instance is destroyed, the pointed-to object is deleted. The move constructor and move assignment operator allow the ownership of an object to be transferred around between `std::unique_ptr` instances (see appendix A, section A.1.1, for more on move semantics). Such a transfer leaves the source object with a `NULL` pointer. This moving of values allows objects of this type to be accepted as function parameters or returned from functions. Where the source object is a temporary, the move is automatic, but where the source is a named value, the transfer must be requested directly by invoking `std::move()`. The following example shows the use of `std::move` to transfer ownership of a dynamic object into a thread:

```
1  void process_big_object(std::unique_ptr<big_object>);
2  std::unique_ptr<big_object> p(new big_object);
3  p->prepare_data(42);
4  std::thread t(process_big_object,std::move(p));
```

By specifying `std::move(p)` in the `std::thread` constructor, the ownership of the `big_object` is transferred first into internal storage for the newly created thread and then into `process_big_object`.

Several of the classes in the Standard Thread Library exhibit the same ownership semantics as `std::unique_ptr`, and `std::thread` is one of them. Though `std::thread` instances don't own a dynamic object in the same way as `std::unique_ptr` does, they do own a resource: each instance is responsible for managing a thread of execution. This ownership can be transferred between instances, because instances of `std::thread` are movable, even though they aren't copyable. This ensures that only one object is associated with a particular thread of execution at any one time while allowing programmers the option of transferring that ownership between objects.

2.3   Transferring ownership of a thread

Suppose you want to write a function that creates a thread to run in the background but passes back ownership of the new thread to the calling function rather than waiting for it to complete, or maybe you want to do the reverse: create a thread and pass ownership in to some function that should wait for it to complete. In either case, you need to transfer ownership from one place to another.

This is where the move support of `std::thread` comes in. As described in the previous section, many resource-owning types in the C++ Standard Library such as `std::ifstream` and `std::unique_ptr` are movable but not copyable, and `std::thread` is one of them. This means that the ownership of a particular thread of execution can be moved between `std::thread` instances, as in the following example. The example shows the creation of two threads of execution and the transfer of ownership of those threads among three `std::thread` instances, `t1`, `t2`, and `t3`:

1

```
1  void some_function();
2  void some_other_function();
3  std::thread t1(some_function);
4  std::thread t2=std::move(t1);
5  t1=std::thread(some_other_function);
6  std::thread t3;
7  t3=std::move(t2);
8  t1=std::move(t3);
```

First, a new thread is started #1 and associated with `t1` . Ownership is then transferred over to `t2` when `t2` is constructed, by invoking `std::move()` to explicitly move ownership #2. At this point, `t1` no longer has an associated thread of execution; the thread running `some_function` is now associated with `t2` .

Then, a new thread is started and associated with a temporary `std::thread` object #3. The subsequent transfer of ownership into `t1` doesn't require a call to `std:: move()` to explicitly move ownership, because the owner is a temporary object—moving from temporaries is automatic and implicit.

`t3` is default constructed #4, which means that it's created without any associated thread of execution. Ownership of the thread currently associated with `t2` is transferred into `t3` #5 , again with an explicit call to `std::move()` , because `t2` is a named object. After all these moves, `t1` is associated with the thread running `some_other_function` , `t2` has no associated thread, and `t3` is associated with the thread running `some_function` .

The final mov e #6 transfers ownership of the thread running `some_function` back to `t1` where it started. But in this case `t1` already had an associated thread (which was running `some_other_function`), so s `td::terminate()` is called to terminate the pr-ogram. This is done for consistency with the `std::thread` destructor. You saw in section 2.1.1 that you must explicitly wait for a thread to complete or detach it before destruction, and the same applies to assignment: you can't just "drop" a thread by assigning a new value to the `std::thread` object that manages it.

The move support in `std::thread` means that ownership can readily be transferred out of a function, as shown in the following listing.

Listing 2.5 Returning a `std::thread` from a function

```
 1  std::thread f()
 2  {
 3      void some_function();
 4      return std::thread(some_function);
 5  }
 6  std::thread g()
 7  {
 8      void some_other_function(int);
 9      std::thread t(some_other_function,42);
10      return t;
11  }
```

Likewise, if ownership should be transferred into a function, it can just accept an instance of `std::thread` by value as one of the parameters, as shown here:

```
 1  void f(std::thread t);
 2  void g()
 3  {
 4      void some_function();
 5      f(std::thread(some_function));
 6      std::thread t(some_function);
 7      f(std::move(t));
 8  }
```

One benefit of the move support of `std::thread` is that you can build on the `thread_guard` class from listing 2.3 and have it actually take ownership of the thread. This avoids any unpleasant consequences should the `thread_guard` object outlive the thread it was referencing, and it also means that no one else can join or detach the thread once ownership has been transferred into the object. Because this would primarily be aimed at ensuring threads are completed before a scope is exited, I named this class `scoped_thread`. The implementation is shown in the following listing, along with a simple example.

Listing 2.6 `scoped_thread` and example usage

```
 1  class scoped_thread
 2  {
 3      std::thread t;
 4  public:
 5      explicit scoped_thread(std::thread t_):
 6          t(std::move(t_))
 7      {
 8          if(!t.joinable())
 9              throw std::logic_error("No thread");
10      }
11      ~scoped_thread()
12      {
13          t.join();
14      }
15      scoped_thread(scoped_thread const&)=delete;
16      scoped_thread& operator=(scoped_thread const&)=delete;
17  };
18  struct func;
19  void f()
20  {
21      int some_local_state;
22      scoped_thread t{std::thread(func(some_local_state))};
23      do_something_in_current_thread();
24  }
```

MEAP

The example is similar to that from listing 2.3, but the new thread is passed in directly to the `scoped_thread` #4 rather than having to create a separate named variable for it. When the initial thread reaches the end of `f #5`, the `scoped_thread` object is destroyed and then joins with #3 the thread supplied to the constructor #1. Whereas with the `thread_guard` class from listing 2.3 the destructor had to check that the thread was still joinable, you can do that in the constructor #2 and throw an exception if it's not.

One of the proposals for C++17 was for a `joining_thread` class that would be similar to `std::thread`, except automatically join in the destructor much like `scoped_thread` does. This didn't get consensus in the committee, so was not accepted into the standard, but is relatively easy to write. One possible implementation is shown in listing 2.7.

Listing 2.7 A `joining_thread` class

```cpp
class joining_thread
        {
            std::thread t;
        public:
            joining_thread() noexcept=default;
            template<typename Callable,typename ... Args>
            explicit joining_thread(Callable&& func,Args&& ... args):
                t(std::forward<Callable>(func),std::forward<Args>(args
            {}
            explicit joining_thread(std::thread t_) noexcept:
                t(std::move(t_))
            {}
            joining_thread(joining_thread&& other) noexcept:
                t(std::move(other.t))
            {}
            joining_thread& operator=(joining_thread&& other) noexcept
            {
                if(joinable())
                    join();
                t=std::move(other.t);
                return *this;
            }
            joining_thread& operator=(std::thread other) noexcept
            {
                if(joinable())
                    join();
                t=std::move(other);
                return *this;
            }
            ~joining_thread() noexcept
            {
                if(joinable())
                    join();
            }
            void swap(joining_thread& other) noexcept
            {
                t.swap(other.t);
            }
            std::thread::id get_id() const noexcept{
                return t.get_id();
            }
            bool joinable() const noexcept
            {
                return t.joinable();
            }
            void join()
            {
                t.join();
            }
            void detach()
            {
                t.detach();
            }
            std::thread& as_thread() noexcept
            {
                return t;
            }
            const std::thread& as_thread() const noexcept
            {
```

```
60                 return t;
61             }
62         };
```

The move support in `std::thread` also allows for containers of `std::thread` objects, if those containers are move aware (like the updated `std::vector<>`). This means that you can write code like that in the following listing, which spawns a number of threads and then waits for them to finish.

Listing 2.8 Spawn some threads and wait for them to finish

```
 1  void do_work(unsigned id);
 2  void f()
 3  {
 4      std::vector<std::thread> threads;
 5      for(unsigned i=0;i<20;++i)
 6      {
 7          threads.push_back(std::thread(do_work,i));          A
 8      }
 9      std::for_each(threads.begin(),threads.end(),
10                    std::mem_fn(&std::thread::join));          B
11  }
```

If the threads are being used to subdivide the work of an algorithm, this is often just what's required; before returning to the caller, all threads must have finished. Of course, the simple structure of listing 2.8 implies that the work done by the threads is self-contained, and the result of their operations is purely the side effects on shared data. If `f()` were to return a value to the caller that depended on the results of the operations performed by these threads, then as written this return value would have to be determined by examining the shared data after the threads had terminated. Alternative schemes for transferring the results of operations between threads are discussed in chapter 4.

Putting `std::thread` objects in a `std::vector` is a step toward automating the management of those threads: rather than creating separate variables for those threads and joining with them directly, they can be treated as a group. You can take this a step further by creating a dynamic number of threads determined at runtime, rather than creating a fixed number as in listing 2.8.

2.4　Choosing the number of threads at runtime

One feature of the C++ Standard Library that helps here is `std::thread:: hardware_ concurrency()`. This function returns an indication of the number of threads that can truly run concurrently for a given execution of a program. On a multicore system it might be the number of CPU cores, for example. This is only a hint, and the function might return 0 if this information is not available, but it can be a useful guide for splitting a task among threads.

Listing 2.9 shows a simple implementation of a parallel version of `std::accumulate`. In real code you'll probably want to use the parallel version of `std::reduce` described in chapter 10, rather than implementing it yourself, but this serves to show some basic ideas. It divides the work among the threads, with a minimum number of elements per thread in order to avoid the overhead of too many threads. Note that this implementation assumes that none of the operations will throw an exception, even though exceptions are possible; the `std::thread` constructor will throw if it can't start a new thread of execution, for example. Handling exceptions in such an algorithm is beyond the scope of this simple example and will be covered in chapter 8.

<div style="border:1px solid">

Listing 2.9 A naïve parallel version of `std::accumulate`

1

2

3

4

5

</div>

```
 1  template<typename Iterator,typename T>
 2  struct accumulate_block
 3  {
 4      void operator()(Iterator first,Iterator last,T& result)
 5      {
 6          result=std::accumulate(first,last,result);
 7      }
 8  };
 9  template<typename Iterator,typename T>
10  T parallel_accumulate(Iterator first,Iterator last,T init)
11  {
12      unsigned long const length=std::distance(first,last);
13      if(!length)
14          return init;
15      unsigned long const min_per_thread=25;
16      unsigned long const max_threads=
17          (length+min_per_thread-1)/min_per_thread;
18      unsigned long const hardware_threads=
19          std::thread::hardware_concurrency();
20      unsigned long const num_threads=
21          std::min(hardware_threads!=0?hardware_threads:2,max_threads);
22      unsigned long const block_size=length/num_threads;
23      std::vector<T> results(num_threads);
24      std::vector<std::thread>  threads(num_threads-1);
25      Iterator block_start=first;
26      for(unsigned long i=0;i<(num_threads-1);++i)
27      {
28          Iterator block_end=block_start;
29          std::advance(block_end,block_size);
30          threads[i]=std::thread(
31              accumulate_block<Iterator,T>(),
32              block_start,block_end,std::ref(results[i]));
33          block_start=block_end;
34      }
35      accumulate_block<Iterator,T>()(
36          block_start,last,results[num_threads-1]);
37      std::for_each(threads.begin(),threads.end(),
38          std::mem_fn(&std::thread::join));              @10
39      return std::accumulate(results.begin(),results.end(),init);
40  }
```

Although this is quite a long function, it's actually straightforward. If the input range is empty #1, you just return the initial value `init`. Otherwise, there's at least one element in the range, so you can divide the number of elements to process by the minimum block size in order to give the maximum number of threads #2. This is to avoid creating 32 threads on a 32-core machine when you have only five values in the range.

The number of threads to run is the minimum of your calculated maximum and the number of hardware threads #3. You don't want to run more threads than the hardware can support

(which is called oversubscription), because the context switching will mean that more threads will decrease the performance. If the call to `std::thread::hardware_concurrency()` returned 0, you'd simply substitute a number of your choice; in this case I've chosen 2. You don't want to run too many threads, because that would slow things down on a single-core machine, but likewise you don't want to run too few, because then you'd be passing up the available concurrency.

The number of entries for each thread to process is the length of the range divided by the number of threads #4. If you're worrying about the case where the number doesn't divide evenly, don't—you'll handle that later.

Now that you know how many threads you have, you can create a `std::vector<T>` for the intermediate results and a `std::vector<std::thread>` for the threads #5. Note that you need to launch one fewer thread than `num_threads`, because you already have one.

Launching the threads is just a simple loop: advance the `block_end` iterator to the end of the current block #6 and launch a new thread to accumulate the results for this block #7. The start of the next block is the end of this one #8.

After you've launched all the threads, this thread can then process the final block #9. This is where you take account of any uneven division: you know the end of the final block must be `last`, and it doesn't matter how many elements are in that block.

Once you've accumulated the results for the last block, you can wait for all the threads you spawned with `std::for_each #10`, as in listing 2.8, and then add up the results with a final call to `std::accumulate #11`.

Before you leave this example, it's worth pointing out that where the addition operator for the type `T` is not associative (such as for `float` or `double`), the results of this `parallel_accumulate` may vary from those of `std::accumulate`, because of the grouping of the range into blocks. Also, the requirements on the iterators are slightly more stringent: they must be at least forward iterators, whereas `std::accumulate` can work with single-pass input iterators, and `T` must be default constructible so that you can create the `results` vector. These sorts of requirement changes are common with parallel algorithms; by their very nature they're different in some manner in order to make them parallel, and this has consequences on the results and requirements. Implementing parallel algorithms is covered in more depth in chapter 8, and the standard supplied ones from

C++17 (the equivalent to the `parallel_accumulate` described here being the parallel form of `std::reduce`). It's also worth noting that because you can't return a value directly from a thread, you must pass in a reference to the relevant entry in the `results` vector. Alternative ways of returning results from threads are addressed through the use of futures in chapter 4.

In this case, all the information required by each thread was passed in when the thread was started, including the location in which to store the result of its calculation. This isn't always the case: sometimes it's necessary to be able to identify the threads in some way for part of the processing. You could pass in an identifying number, such as the value of `i` in listing 2.8, but if the function that needs the identifier is several levels deep in the call stack and could be called from any thread, it's inconvenient to have to do it that way. When we were designing the C++ Thread Library we foresaw this need, and so each thread has a unique identifier.

2.5   Identifying threads

Thread identifiers are of type `std::thread::id` and can be retrieved in two ways. First, the identifier for a thread can be obtained from its associated `std::thread` object by calling the `get_id()` member function. If the `std::thread` object doesn't have an associated thread of execution, the call to `get_id()` returns a default-constructed `std::thread::id` object, which indicates "not any thread." Alternatively, the identifier for the current thread can be obtained by calling `std::this_thread:: get_id()`, which is also defined in the `<thread>` header.

Objects of type `std::thread::id` can be freely copied and compared; they wouldn't be of much use as identifiers otherwise. If two objects of type `std::thread::id` are equal, they represent the same thread, or both are holding the "not any thread" value. If two objects aren't equal, they represent different threads, or one represents a thread and the other is holding the "not any thread" value.

The Thread Library doesn't limit you to checking whether thread identifiers are the same or not; objects of type `std::thread::id` offer the complete set of comparison operators, which provide a total ordering for all distinct values. This allows them to be used as keys in associative containers, or sorted, or compared in any other way that you as a programmer may see fit. The comparison operators provide a total order for all non-equal values of `std::thread::id`, so they behave as you'd intuitively expect: if `a<b` and `b<c`, then `a<c`, and so forth. The Standard Library also provides

`std::hash<std::thread::id>` so that values of type `std::thread::id` can be used as keys in the new unordered associative containers too.

Instances of `std::thread::id` are often used to check whether a thread needs to perform some operation. For example, if threads are used to divide work as in listing 2.9, the initial thread that launched the others might need to perform its work slightly differently in the middle of the algorithm. In this case it could store the result of `std::this_thread::get_id()` before launching the other threads, and then the core part of the algorithm (which is common to all threads) could check its own thread ID against the stored value:

```
1  std::thread::id master_thread;
2  void some_core_part_of_algorithm()
3  {
4      if(std::this_thread::get_id()==master_thread)
5      {
6          do_master_thread_work();
7      }
8      do_common_work();
9  }
```

Alternatively, the `std::thread::id` of the current thread could be stored in a data structure as part of an operation. Later operations on that same data structure could then check the stored ID against the ID of the thread performing the operation to determine what operations are permitted/required.

Similarly, thread IDs could be used as keys into associative containers where specific data needs to be associated with a thread and alternative mechanisms such as thread-local storage aren't appropriate. Such a container could, for example, be used by a controlling thread to store information about each of the threads under its control or for passing information between threads.

The idea is that `std::thread::id` will suffice as a generic identifier for a thread in most circumstances; it's only if the identifier has semantic meaning associated with it (such as being an index into an array) that alternatives should be necessary. You can even write out an instance of `std::thread::id` to an output stream such as `std::cout`:

```
1  std::cout<<std::this_thread::get_id();
```

The exact output you get is strictly implementation dependent; the only guarantee given by
the standard is that thread IDs that compare as equal should produce the same output, and
those that are not equal should give different output. This is therefore primarily useful for
debugging and logging, but the values have no semantic meaning, so there's not much more
that could be said anyway.

2.6   Summary

In this chapter I covered the basics of thread management with the C++ Standard Library:
starting threads, waiting for them to finish, and not waiting for them to finish because you
want them to run in the background. You also saw how to pass arguments into the thread
function when a thread is started, how to transfer the responsibility for managing a thread
from one part of the code to another, and how groups of threads can be used to divide work.
Finally, I discussed identifying threads in order to associate data or behavior with specific
threads that's inconvenient to associate through alternative means. Although you can do
quite a lot with purely independent threads that each operate on separate data, as in listing
2.9 for example, sometimes it's desirable to share data among threads while they're
running. Chapter 3 discusses the issues surrounding sharing data directly among threads,
while chapter 4 covers more general issues surrounding synchronizing operations with and
without shared data.

MEAP

Up next...
3 Sharing data between threads
- Problems with sharing data between threads
- Protecting data with mutexes
- Alternative facilities for protecting shared data

© 2018 Manning Publications Co.